Multilingual NLP

HeRE: A Novel Benchmark for Hebrew Relation Extraction via LLM-Guided Denoising of Knowledge Graph Alignments [arxiv]

Ron Keinan, Amir David Nissan Cohen, Reut Tsarfaty

Build a large native-Hebrew RE benchmark by treating noisy KG–sentence matches as an NLI entailment check and using LLMs to denoise at scale—no translation, little manual labeling.

HeRE: a Hebrew relation-extraction benchmark from weak KG-to-text alignments, with the key idea of validating each candidate triple by asking whether the sentence entails a Hebrew verbalization of the triple (NLI framing), enabling ~1.5M silver instances while avoiding English-centric artifacts. Mine candidate KG triple–sentence pairs where both entities appear; convert each triple to a Hebrew hypothesis and run entailment filtering (premise=sentence). Compare 5 denoising strategies across few-shot prompting, NLI fine-tuning, and label-free cross-training; test Hebrew vs English prompting; evaluate multilingual vs Hebrew-specific models. Multilingual LLMs lead in few-shot prompting, but Hebrew-specific encoders become strongest with NLI fine-tuning and in label-free cross-training. Hebrew prompts match or slightly beat English prompts. Best open model (Gemma-4-31B-it) reaches F10.915 vs best API (GPT-5.4) F10.921; pipeline scores F10.844 on a fully held-out dual-annotated test set.

Disentangling Language Effects on Tool Admission and Post-Retrieval Safety [arxiv]

Aditya Nawal, Manit Baser, Mohan Gurusamy

Changing the language of a harmful request can sharply increase whether agents choose to retrieve risky URLs, so multilingual safety must separate "tool admission" from "post-retrieval harm."

Separates two often-confounded safety failures: (1) tool admission (does the agent retrieve?) and (2) post-retrieval safety (is the final answer harmful). Introduces SLIP (Safety Loss through Interlingual Processing): safety degrades systematically across languages, and "non-harmful" outputs can mask refusal vs deflection vs capability failure. Controlled URL-retrieval testbed with independently set languages for request, retrieved page, and required answer. Evaluate 10 models across 7 languages; measure admission rate on harmful requests and harmfulness conditional on matched tool calls; replicate in a more autonomous retrieval setting; audit "safe" outputs into refusal/deflection/unusable categories. Non-English harmful requests raise tool admission from 46.9% to 67.4% on average (increase for every model). For matched tool calls, forcing answers in the page language increases pooled harmfulness (50.3%57.4%) for several closed models, especially in Chinese/Hindi; translating pages alone shows no consistent increase. "Non-harmful" labels often hide qualitatively different behaviors, so single-score safety evaluations can misdiagnose failures.

On the Tip of the Tongue: First-Token Cueing as a Behavioral Probe of Recall Failures in LLMs [arxiv]

Jaewon Sok, Jewon Yeom, Seonghyeon Park, Jeongjae Park, Hwiyeong Lee, Taesup Kim

Many LLM factual errors are first-step decoding failures: forcing just the gold first token often unlocks the correct multi-token answer, across languages.

First-Token Cueing as a behavioral probe: prefill the model's response with the first token of the correct answer to test "tip-of-the-tongue" failures. Argues a sizable share of QA hallucinations reflect recoverable retrieval/decoding bottlenecks at the first token, not missing knowledge. Compare standard closed-book QA generation vs generation constrained to start with the gold answer's first token (and, in analysis, longer gold prefixes). Measure accuracy recovery on multi-token answers; analyze uncertainty/confidence around the first token; test whether models more confidently continue correct starts and whether alternative sampled paths can solve the same items. Across 16 open models, a single gold first token fixes ~20–50% of multi-token answer errors; the first token yields the largest recovery among all prefix lengths, localizing the failure to the initial decoding decision. Gains hold on a 7-language multilingual benchmark, and models often continue the correct start with higher confidence despite rarely finding it unaided.

KiranaBench: Can a Language Model Run a Neighbourhood Shop? [arxiv]

Bhuvan Arora, Ankita Karmakar, Chaitanya Jain, Dilesh Churiwal, Harshit Jain, Shivansh Manish Dubey, Manish Prasad, Raja Piyush, Kshitij Tayal, Dhruv Kumar, Jagat Sesh Challa, Raunak Manekar, Yash Sinha

In a realistic multilingual, multi-role Indian kirana simulation, open LLM agents can underperform simple rules and sometimes make the shop worse than doing nothing.

KiranaBench: long-horizon neighbourhood-shop simulation emphasizing operational retail work (inventory, customers, cash, bookkeeping) with four separate LLM roles and interactions in English/Hindi/Hinglish, using India-calibrated demand—exposing coordination and multilingual brittleness missed by prior agent benchmarks. Simulate manager/cashier/helper/bookkeeper as separate agents with partial information that must communicate. Generate supplier/customer interactions in English, Hindi, Hinglish; demand grounded in India's Household Consumption Expenditure Survey. Evaluate open-weight LLMs vs a fixed-rule baseline using business outcomes (profit; whether performance is worse than unmanaged). All tested open models fail to run the shop competitively: best earns ~1/8 the profit of the fixed-rule baseline; two models perform so poorly the shop ends worse than unmanaged. Performance drops when supplier communication is in Hindi/Hinglish vs English, indicating a real multilingual reliability cliff in long-horizon multi-agent settings.

Steering beyond Language: Improving Multilingual Training-Free Text Embeddings via Paired-Residual Subspace Removal [arxiv]

Ailiang Lin, Zhuoyun Li, Yusong Wang, Keyu Mao, Kotaro Funakoshi, Manabu Okumura

Frozen LLM embeddings are skewed by a low-rank "language identity" component; subtracting a residual subspace learned from parallel text improves multilingual alignment without training.

Paired-Residual Subspace Removal (PRSR): a training-free inference-time fix that treats language effects as a removable low-dimensional nuisance factor. Uses translation pairs to estimate directions that change with language (not meaning) and projects them out to make embeddings more language-invariant. From parallel sentence pairs, compute residual vectors capturing representation shifts across languages; fit a low-rank subspace (layer-wise) summarizing language-induced variation. At inference, take the chosen embedding (e.g., last-token hidden state) and remove its projection onto this residual subspace—no weight updates. PRSR yields consistent gains over training-free embedding baselines across multilingual retrieval evaluations (MTEB/MMTEB) covering 18 languages and cross-lingual bitext mining with English paired to 200+ languages, supporting the claim that much multilingual embedding weakness in frozen LLMs comes from a cheap-to-remove low-dimensional language component.

Multilingual Speech Processing

SEA-HELM Audio: Evaluating Multilingual ASR-LLM Pipeline Performance Across SEA-HELM Subsets [arxiv]

Aulia Adila, Yixian Tan, Jonathan Heng, Kunat Pipatanakul, Donghang Wu, Anocha Sutaveephamochanon

Text-strong models can fail badly as voice agents; SEA-HELM Audio quantifies the end-to-end "speech gap" across Southeast Asian languages and tasks.

SEA-HELM Audio adds a speech layer to SEA-HELM, showing ASR is not a neutral front-end: transcription errors interact with downstream reasoning/instruction following. Introduces the "speech gap" (text accuracy speech accuracy) as a comparable metric for multilingual voice-agent reliability. Run SEA-HELM tasks in parallel: (1) original text prompts to the LLM, (2) prompts synthesized to speech ASR transcript same LLM. Keep scoring/verifiers identical so differences isolate ASR+downstream effects. Evaluate open and commercial ASR-LLM pipelines across Chinese, English, Malay, Indonesian, Tamil. Speech input causes large drops: multiple-choice speech gap 9.00–18.50 points; instruction-following gap 20.00–42.86 points. Tamil is most fragile, with >40-point degradation, highlighting language equity issues and showing text-only benchmarks poorly predict spoken performance.

Anuvani-LID: Robust and Calibrated Spoken Language Identification for Indian Languages [arxiv]

Stephen Fernandes, Narayan Choudhary

Spoken language ID is not "free" for Indian languages in general speech/audio-LLM systems; a dedicated Conformer model (Anuvani-LID) is far more accurate and well-calibrated.

Demonstrates concrete LID gaps in general-purpose models (limited language-token coverage; poor short-clip performance) and releases Anuvani-LID: a robust, calibrated Indian spoken LID model evaluated under many controlled real-world conditions, not just clean accuracy. Train a 115M-parameter Conformer classifier on 54,643 hours spanning 22 Indian languages + Indian-accented English.

Evaluate on a 102-condition controlled benchmark with speaker-held-out and corpus-held-out splits; stress short/noisy/narrowband/code-mixed audio; report calibration via ECE for routing reliability. General systems underperform on 1-second clips (Gemini 3.5 Flash macro-F1=0.52; GPT-audio=0.14) and Whisper-like models expose LID tokens for only 15/23 languages. Anuvani-LID reaches 0.914 mean macro-F1 over 37 conditions (5.2 points above Gemini), is +24 points over Gemini at 1 second, scores 0.897 macro-F1 on all 23 languages (no other system >0.713), and is tightly calibrated (ECE=0.008).

Field Overview

Across the list, the dominant theme is LLM/agent reliability, safety, and evaluation: at least ~20 papers focus on prompt injection/jailbreaks, refusal/harms, watermarking/detection, PII/de-identification, and uncertainty/hallucination measurement (e.g., multiple works on tool-use safety, RAG hallucination risk, watermarkability, PII detection, and jailbreak backdoors). A second major cluster is agentic systems engineering—roughly ~20 papers on long-horizon agents, memory (filesystem/secure/governed memory, memory selection/gating), tool discovery/recommendation, routing/cascades, and new benchmarks/auditing frameworks (RealCLI, KiranaBench, AuditKIT, SIT-Eval, Annotrace, etc.). Efficiency and inference/training optimization is also heavily represented (15+): speculative decoding variants, MoE routing/expert reduction, width slicing/elastic models, embedding/tokenizer alternatives, quantization/low-rank optimizers, and attention variants. A notable emerging direction is multimodal/audio + clinical integration (10–15 papers) spanning medical audio for LLMs, speech emotion/judging, TTS/one-step generation, and robust ASR/deepfake/spoof detection, alongside a steady stream of multilingual/low-resource language work (8–12) including dialectal Arabic, Luganda, Hebrew RE, Bengali idioms/hallucination detection, and cross-lingual safety/tooling.