Multilingual NLP

DeepSEA: A Multilingual Evaluation Benchmark for Southeast Asia [arxiv]

Jiuzhou Han, Abhishek Bijay Mishra, Trevor Cohn, Shubham Mittal

DeepSEA benchmarks 12 Southeast Asian languages and shows many "multilingual" models reason better by pivoting through English than directly in-language.

DeepSEA: professionally translated, task-diverse evaluation (reasoning, QA, summarization, MT) for SEA languages, enabling diagnosis of true in-language capability vs translation dependence. Human-translate established English datasets into 12 SEA languages; evaluate proprietary and open models; test "English pivoting" (solve via English) and analyze cross-lingual errors plus tokenizer fertility. Big resource-tier gap: frontier models stay more robust, while smaller models degrade sharply on Burmese/Khmer/Lao. For reasoning/summarization, many models improve with English pivoting, revealing English-anchored "thinking"; tokenization/script issues further explain failures.

One Head Away from Language Confusion: Single Attention Heads Across Model Families [arxiv]

Jiwon Jeong, Anjelo Jann Laroza, Chaewoon Yoon, Arjun Pillai, Pranav Kokati, Christian Hoang, Seunghyeok Hong

Wrong-language replies can hinge on a single attention head; ablating or replacing it can reliably induce or steer language switching.

Causal localization of language control: across many instruction-tuned multilingual models, one dominant head disproportionately enforces the requested output language, acting like a controllable switch. Head-by-head ablations on FLORES-style continuations and the Language Confusion Benchmark; validate with mean ablation and mean-replacement at language-name tokens; steer by replacing the head's output with another language's mean. 10/12 models have a single head whose removal uniquely spikes language confusion; in 7 models, ablation shifts ~8%–51.5% of continuations out of-language (others ~1.4%). Mean-replacement redirects 84%–99% of non-English LCB replies into a chosen language in 5 models; effect is weaker in base (non-instruction) models.

Training with Ciphers to Enhance In-Context Terminology Learning For Translation from Rare Languages [arxiv]

Deevyanshu Malu, Deeptanshu Malu, Sunita Sarawagi

Cipher-based fine-tuning forces translation models to learn rare term translations from in-context examples, boosting terminology recall without hurting overall quality.

Defines in-context term learning (ICTL) as a missing capability and proposes RareCipherFT, a resource-free recipe (no dictionaries/term lists) to strengthen rare-term uptake from provided context while preserving general translation. During in-context fine-tuning, choose context examples sharing a rare term with the target, replace that term with a consistent character-shuffled cipher in both context and reference, and anneal the ciphered fraction over training. Across 7 directions (EnglishTamil/Assamese/Mizo/Nepali/Hausa) and 1B–7B models, RareCipherFT improves rare-term recall by ~5–15 absolute points vs standard IC-Train with no COMET drop; controls suggest gains come from true context-based term linking, not shortcuts.

Efficiency Measured in English: Tokeniser Coverage Bounds Small Language Models on Indic Scripts [arxiv]

Anna Pramod

Small-model "efficiency" can be English-only: poor tokenizer coverage on Indic scripts forces byte fallback, exploding token counts and shrinking usable context.

Identifies tokenizer vocabulary coverage as a binding constraint for multilingual SLMs; shows byte fallback can hide catastrophic script failure (0% UNK yet mostly bytes), so efficiency claims must include tokenization diagnostics. On a 131k-sentence, 11-script Indic corpus, measure fertility (tokens/word), codepoint coverage, and UTF-8 byte-fallback rates; compare tokenizers/models (e.g., TinyLlama, mBERT, MuRIL). TinyLlama tokenization jumps from ~1.34 tokens/word (Latin) to 22.60 (Kannada) due to near-zero script coverage (0% Odia, 1% Telugu, 2% Kannada) and heavy byte fallback (e.g., Odia at 99.6% of byte ceiling; 93.6% tokens are bytes) while still reporting 0% UNK. Practical hit: a 2048-token window fits ~1,528 English words but ~91 Kannada words; MuRIL

shows better Indic coverage can fix this with fewer parameters.

Multilingual Speech Processing

TeleASR-Amharic: A Benchmark for Robust Amharic ASR in Realistic Telecommunication Conditions [arxiv]

Dagnachew Mekonnen Marilign, Enrique Molero, Abrham Haile, Louis Powell, Israel Abebe Azime

They release a telecom-realistic Amharic ASR benchmark and show clean-audio gains don't translate: WER can collapse under channel/noise shifts.

TeleASR-Amharic: a reproducible evaluation suite capturing real telecommunication degradations (channel + environment) to measure robustness, exposing that low-resource ASR progress is overstated on overly clean test sets. Analyze degradations in real telecom recordings, build multi-condition benchmark, and test robustness via training-data scaling and telecom-aware augmentation; evaluate with WER across severity levels. Severe telecom conditions can make a strong clean model unusable (WER 0.230.79; +243% relative errors). More data plus telecom-aware augmentation recovers robustness, with up to 23.6% relative WER reduction under extreme noise.

Monsoon: Scalable Spontaneous Speech Collection for Low-Resource ASR in 50 Languages [arxiv]

Shobhit Banga, Manas Dhir, Aaditya Pareek, Amritansh Walecha, Bhaskar Singh, Hanuman Sidh, Manmeet Kaur, Sagar Jain, Supriya Paul, Vanshika Chhabra, Utkarsh Singh

! Summary may contain inaccuracies: The summary states they "do not claim to release the collected corpora themselves." The abstract says they "release the framework, its quality-control models, and the fine-tuned models," but it does not state whether the collected corpora are released or not. Claiming they do not release the corpora adds information not supported by the abstract.

Monsoon argues data—not architectures—is the ASR bottleneck, and provides a scalable pipeline to collect high-quality spontaneous speech in 50 languages with large conversational gains.

Monsoon: end-to-end framework for eliciting, filtering, and multi-stage correcting spontaneous conversational speech for low-resource ASR; releases the framework, QC models, and fine-tuned ASR models (not necessarily the raw corpora). Pair contributors to elicit unscripted dialogue; automated screening for low-quality/problematic audio; five-stage human transcript review; coverage-aware segment selection to broaden lexical/acoustic diversity. Collected across 50 languages; analysis on seven ~3,500-hour corpora shows the staged QC reduces semantic WER 10–39% at fixed hours vs conventional collection. Whisper fine-tuned on Monsoon matches a public-corpora baseline on that baseline's tests but improves 20–44 semantic WER points on conversational benchmarks; gains transfer to another ASR architecture and scale with more Monsoon data.

Foundation Model Theory

Cognitive Map in Large Language Models [arxiv]

Albert Qin, Jonathan Cohen, Gautam Reddy

LLMs encode relations via a separable "structure + entity binding" representation that can be causally swapped or edited to change inferred relationships.

Turns "cognitive map" from metaphor into mechanism: relational structure is represented independently of entity identities and combined via role–filler binding, enabling predictable activation splicing and targeted memory edits. Relational inference tasks (incl. novel entity tokens) plus causal "mosaic" activation splicing to swap structure vs identity; identify sparse key–value-like binding components and perturb them to rewrite stored links. Large (~10B+) open models generalize relational inference to new entities; mosaics yield hybrid answers consistent with factorized structure/identity. Editing the sparse binding substrate can implant specific false relational memories with targeted downstream effects.

LLM Evaluation

Rubrics Carry the Standard: Shapley Attribution for Checklist-Based LLM Judges [arxiv]

Saeid Asgari, Rui Ying, Sara Malvar, Rakshanda Agarwal, Bruce Sun, Claire Wang, Leonardo de Oliveira Nunes, Ranveer Chandra

Checklist rubrics largely set LLM-judge strictness, and LLM "self-review" of rubrics drifts standards; Shapley attribution can detect and block harmful edits.

Shows the rubric—not the judge model—dominates outcomes, and introduces exact, linear-time Shapley attribution for all-must-pass checklists to quantify each check's unique value vs false-fail harm; adds change-control and provenance freeze to prevent drift. Evaluate rubric sensitivity on a planted-flaw grading task plus SummEdits and RAGTruth; compute closed-form Shapley per check; use these signals to gate rubric revisions and prevent deletion of human-derived checks. With judge/output fixed, rubric edits swing correct-output pass rates from 2% to 100%, while swapping judge models or re-rendering checks changes only a few points. Self-review drifts across 3 revisers and all datasets, often deleting uniquely useful checks; the Shapley-based gate + provenance freeze blocks nearly all degrading revision chains and yields identical decisions across three judge models.

Field Overview

Across today's 1556 papers, the dominant theme is LLM agents and their infrastructure: at least ~120 titles focus on agent memory (persistent/latent memory, context compression, KV-cache reuse), tool use, long-horizon evaluation harnesses, and multi-agent coordination (e.g., many "Bench/Env/Harness" papers plus memory governance and trace auditing). A second major cluster is reasoning/post-training mechanics—roughly ~70 papers on RLVR/GRPO/OPD/self-distillation, test-time scaling/early stopping, rubric/judge reliability, and diffusion language models/parallel decoding/speculative decoding (often tied to efficiency and cost). Multilingual/low-resource + safety is another noticeable concentration (~40+), with multiple Bangla/African/Arabic/SEA benchmarks and cross-lingual safety defenses, tokenizer coverage issues, and culturally aligned MT. A surprising emerging direction is "System-1 / decision models" (Jev-style typed decision models, calibration under shift, cost-aware routing) appearing repeatedly alongside agent governance, suggesting a shift toward probabilistic, auditable decision components rather than pure text generation.