Multilingual NLP

AQL-Lite: Enabling Efficient Knowledge Acquisition for Arabic LLMs via Data-Centric Continual Pretraining [arxiv]

Houssam Eddine Boukhalfa, Mourad Bouache, Mourad Dahmane

Arabic gains come from curated Modern Standard Arabic data plus partial-parameter continual pretraining—not raw web text or full retraining.

AQL-Lite: a practical, data-centric MSA curation pipeline paired with an efficient continual-pretraining recipe that updates only selected weights, delivering most of full-parameter gains at much lower cost. Build an MSA-focused "AQL-curated" corpus via cleaning/filtering, MSA selection, and deduplication; continual pretrain with RS-LoRA adapters plus unfreezing embeddings and last 8 transformer blocks using layer-wise learning rates. At 1B tokens, raw-web Arabic yields ~no improvement, while AQL-curated improves mean over 4 Arabic benchmarks by +3.51; at 5B tokens, partial-parameter training reaches 94% of full-parameter gain (73.19 vs 73.58) in 52% wall-clock; scaled to 40B MSA + 1M instruction-tune hits 75.3 on OALL v2 (7 benchmarks) for a 12B model.

Same Entity, Different Script: Diagnosing Cross-Script Entity Matching Failures in Language Models [arxiv]

Athira Gopal, Patil Prathamesh Pravin, Ankush Kadu, Arpit Jain, Ashwanth Krishnan

Cross-script entity matching fails mainly due to a miscalibrated decision threshold, not missing knowledge—and a targeted activation intervention fixes much of it.

Diagnoses a script-specific "same person?" failure where internal representations still separate matches from non-matches, but the final criterion is too conservative in native scripts; localizes a causal site and derives a fact-independent intervention direction. Create audited Latinnative-script identity-pair datasets with controlled templates; evaluate same/different asymmetry; train probes on Latin hidden states and transfer to native-script states; show scalar calibration explains much of the gap; use activation patching in Qwen2.5-7B to find where entity-2 info affects the decision and apply a sensitivity-direction edit. Qwen2.5-7B (Marathi): same-person accuracy 35.1% (Latin) vs 8.5% (Devanagari) while different-person stays >98%; Malayalam set: same-person 5.3% but AUROC 0.925; calibration explains 87% of probe–model gap; intervention boosts held-out Marathi same-person 6.5%33.5% with different-person 98.5%.

CrisiLangs: A Multilingual Resource for Crisis Warning Generation [arxiv]

Giacomo Gonella, Stefano Menini, Marco Guerini

CrisiLangs provides expert-aligned, parallel multilingual crisis-warning data and shows fine-tuning beats zero-shot for reliable public alerts.

Extends English CrisiText into a controlled multilingual benchmark (Italian, Bulgarian, German, Greek) with expert translations aligned to professional warning guidelines, enabling apples-to-apples multilingual generation studies. Translate 1,000 crisis events' guideline-based English warnings into 4 languages (5,000 aligned pairs total); fine-tune a multilingual model jointly; test prompting with vs without guidelines; evaluate with ROUGE-L/BLEU/BERTScore plus human preferences across decoding temperatures. Multilingual fine-tuning consistently outperforms zero-shot across all languages/metrics; humans prefer low-temperature (more deterministic) warnings; guideline-in-prompt effects vary by language, suggesting prompt transfer is not uniform.

Measuring Research Progress and Linguistic Inclusion in Low-Resource Languages [arxiv]

Atnafu Lambebo Tonja, Tadesse Destaw Belay, Jesujoba Oluwadara Alabi, Israel Abebe Azime, Benjamin Rosman, Ehsan Shareghi, Thamar Solorio, Pontus Stenetorp, Monojit Choudhury

Low-resource NLP output is growing fast, but linguistic inclusion is not—this paper introduces metrics to measure whether coverage is actually becoming equitable.

Builds a 2020–2025 dataset of 1,158 low-resource/multilingual NLP papers and proposes the Linguistic Inclusion Framework (LIF) with four measurable dimensions to separate "more papers" from "broader, sustained, equitable language coverage." Extract languages per paper and aggregate over time/region; compute LIF: (1) research growth, (2) representation expansion (new languages), (3) representation continuity (repeat attention), (4) distributional equity (representation share vs speaker-population share) at continent/country levels.

Papers rise 52415/year (5.4Œ share within ACL Anthology), but growth concentrates on already-studied languages; Asia has ~64.1% of speakers but only 46.7% representation (~0.7Œ), while Europe gets ~5.3Œ its proportional share; sub-Saharan Africa and Southeast Asia are especially undercovered, and no continent scores well across all LIF dimensions.

Multilingual Speech Processing

Group Relative Policy Optimisation Improves Multilingual Speech Recognition in Low-Resource Languages [arxiv]

Tim Wenjie Hu, Evelyn Nafula Ouma, Benjamin Akera, John A. Quinn

GRPO RL fine-tuning reliably improves multilingual Whisper on 51 low-resource African languages, especially under noise and scarce data.

Shows GRPO transfers from high-resource/LLM-style settings to a standard encoder–decoder ASR model in a truly low-resource multilingual regime, mitigating exposure bias by training on the model's own rollouts; releases SunflowerASR. Start from Whisper-large-v3, add a GRPO stage: sample a small group of candidate transcripts per utterance, score with a sequence-level ASR reward (WER-aligned), and increase probability of better candidates relative to worse ones (no separate value baseline), encouraging recovery after early errors. Best cited macro-WER on African benchmarks: 0.227 on AfriVox-v2 and 0.372 on SimbaBench; gains are largest with limited data and degraded/noisy audio, matching the exposure-bias hypothesis.

Synthetic-to-Real Code-Switched ASR: A Spontaneous Nepali-English Podcast Benchmark [arxiv]

Firoj Paudel

A new audited Nepali–English code-switched podcast benchmark shows synthetic fine-tuning drives most gains; RL adds little once SFT converges.

Releases a real-world spontaneous code-switched ASR benchmark with contamination audits and diagnostics (switch-point errors, script confusion), and tests a low-cost synthetic-to-real adaptation recipe when labeled code-switched audio is scarce. Stage 1: generate ~0.95h multi-speaker code-switched TTS audio and do supervised fine-tuning; Stage 2: optional GRPO with group-sampled transcripts and rewards for transcription quality and optionally script correctness; evaluate with MER, SPER, and script confusion rates. Synthetic SFT reduces MER 23.40%21.52% and SPER 35.89%28.77% on the podcast test set; GRPO yields only marginal changes (e.g., English script confusion 6.45%6.36%, SPER 28.77%28.70%); explicit script reward doesn't help because GRPO samples show near-zero script variation post-SFT; Nepali improves while English is largely preserved on monolingual tests.

Foundation Model Theory

Verbalizing Activation Differences with Natural Language Autoencoders [arxiv]

Andrea Blasi Núñez, Gianluca Barmina, Federico Torrielli, Thiusius Rajeeth Savarimuthu, Peter Schneider-Kamp, Lukas Galke Poech

Natural language autoencoders can "read" activation differences, enabling reusable model-diff verbalizations that often reveal finetuning objectives or hidden trigger words.

Demonstrates NLAs generalize beyond their exact training setup and extends them to verbalize activation *differences* between base and finetuned models, turning them into a human-readable internal change detector. Test (1) compositionality via activation arithmetic/analogies, (2) reuse of an NLA across lightweight finetunes without retraining, (3) diffing by feeding activation deltas from "Taboo model organisms" vs base to recover secret words/objectives, and (4) RL/specialized training on differences to improve held-out diffing. Verbalizations behave compositionally (often naming the intended target and not the subtracted concept); an NLA transfers to related finetunes with modest degradation (largest model drops more); diff verbalization frequently recovers concealed secret words, and RL on differences improves recovery on held-out organisms; informative positions can occur pre-prompt, though NLAs still trail specialized misalignment-diff methods while beating black-box baselines.

Field Overview

Today's batch is dominated by LLM evaluation/reliability and agentic systems: at least ~25–30 papers focus on benchmarks and measurement issues (LLM-as-a-judge robustness, reward-score comparability, fact-checker evaluation, hallucinated citations, prompt sensitivity, open-book/RAG scoring, and domain benchmarks like IFC-MCQ, TutorCaseBench, TRIAD, IndoNutriQA). A second major cluster is long-horizon agents with memory and policy control (at least ~12–15 on memory-conditioned planning, preference acquisition, personal-agent memory benchmarks, post-deletion persistence, confidentiality "mosaic" leaks, and compiling natural-language policies/guardrails). Multilingual/low-resource inclusion is also unusually prominent (at least ~10–12 on Arabic/Bashkir/Vietnamese/Slovak/Nepali-English code-switching, cross-script entity matching, crisis warnings, and linguistic inclusion/bias), alongside a smaller but clear efficiency/architecture thread (~6–8 on KV-cache compression, low-bit n-gram embeddings/VQ training, model fusion, sparse MoE, and recurrent/looped models). A notable emerging direction is "safety as systems engineering" for agents—papers connect guardrails, privacy leakage, and memory persistence into end-to-end risk models rather than single-turn moderation.