Multilingual NLP

A Russian, a Brazilian, and a German Walk into a Bar, Ask the Same Question– And Get Three Different Answers: Profiling LLM Perspective Diversity on Contested Topics [arxiv]

Wiem Ben Rim, Yao Lu, Pontus Stenetorp, Emine Yilmaz

LLMs can shift stance on contested questions across languages; the paper introduces a dataset and framework to measure this cross-lingual "perspective drift."

Introduces PerspectivesQA: manually verified contested questions in 18 languages, and ROOTS: a profiling framework that disentangles language effects, stance/perspective variation, and framing/surface-form sensitivity to quantify worldview changes beyond English-only evals. Curate multilingual contested questions with manual verification; profile eight open/proprietary LLMs with ROOTS by varying query language, framing, and small prompt perturbations, then measure stance/commitment and sensitivity across conditions. Across eight models, answers hedge more in English but commit more in lower-resource languages; larger models are more cross-lingually consistent than smaller ones, yet all models show notable stance shifts from minor surface-level prompt changes.

When the Weakest Language Keeps Moving: ARMOR-DP for Multilingual BabyLM [arxiv]

Hanyu Liao, Yijing Chen, Tao Ji

In multilingual training, the "weakest language" changes over time; dynamic sampling based on current per-language loss beats static protection of a pre-chosen language.

ARMOR-DP: a diagnose-then-control framework that identifies why a language is lagging (tokenization, exposure, capacity, objective) and adjusts sampling using per-language validation loss so protection follows the moving bottleneck. Wrap BabyLM multilingual training in a control loop: monitor per-language validation loss, update sampling to upweight the current worst language, and separately test/attribute imbalance sources (representation artifacts, exposure, capacity, objective pressure). Weakness can be an artifact (e.g., equal character budgets hurt Chinese; byte-level encoding reduces this), and once one language is protected another (e.g., English) can become weakest. ARMOR-DP tracks these shifts, improves balance vs static "protect Chinese," and matches/beat a static baseline with fewer Chinese epochs; scaling capacity further helps.

Merge-Based Vocabulary Extension Stalls at the Pre-Tokenizer: An English-Preserving Repair for Nepali, Tested in Continued Pretraining [arxiv]

Aarjan Chaudhary (Naamche Labs)

Nepali tokenization in several LLMs is bottlenecked by a regex pre-tokenizer (not BPE merges); a tiny backward-compatible regex fix enables near-English token efficiency.

Identifies a pre-tokenizer bug: "word" regex includes Unicode letters but excludes combining marks, so Devanagari vowel signs split words before BPE. Proposes a minimal regex repair that preserves English token IDs for text without Devanagari/combining marks; releases repaired tokenizers and code. Analyze the regexBPE pipeline and show BPE cannot merge across regex-created boundaries; modify regex to keep Devanagari combining marks within word spans; evaluate token counts and run continued pretraining on Llama-3.2-1B and Qwen3-0.6B. Broken pre-tokenization imposes a ~2.2Œ Nepali token "floor" vs English even with unlimited merges; many existing Nepali extensions already hit this ceiling. With the regex fix + 32k merges, Nepali drops to ~1.01Œ English token usage (~54% fewer tokens vs standard extension). Continued-pretraining gains are recipe/compute-dependent: token savings don't automatically improve held-out Nepali bits-per-byte without controlling steps/data exposure.

Multilingual Speech Processing

Preserve to Expand: Source-Aware Adaptation and Model Merging for Multilingual LLM-Based ASR [arxiv]

Yuchen Zhang, Haralambos Mouratidis, Ravi Shekhar

New ASR languages can be added to a frozen-LLM Speech LLM without source-data replay by building source-aware connector updates and merging them to reduce forgetting.

Systematic study of connector-only adaptation for multilingual ASR under a no-source-data constraint, showing forgetting depends on update geometry and how multiple language updates are merged (not just training amount). Use a connector-based Speech LLM (speech encoder trainable connector frozen LLM). Compare target-language connector update construction (naive target FT vs source-aware/importance-aware like SSU-Wanda) and update combination (e.g., TIES; multi-target merge vs sequential/recursive expansion). Naive target-only connector tuning (and straightforward TIES) often learns the target but severely degrades source languages. Best single-target trade-off: SSU-Wanda + TIES. For multiple targets, merging several target updates is usually more source-stable than sequential expansion, though ordering in recursive expansion can improve target performance.

SimPOTrans: Self-Adaptation for Offline-to-Simultaneous Speech Translation via Reinforcement Learning [arxiv]

Haotian Tan, Hiroki Ouchi, Sakriani Sakti

A simple RL post-training turns an offline speech translation model into a strong simultaneous translator without explicit read/write policies; latency is controlled by chunk size at test time.

SimPOTrans: self-adaptation via RL that rewards useful partial-context translations, avoiding synthetic read/write trajectories and explicit latency losses while still enabling quality–latency trade-offs at inference. Reframe SimulST as multi-turn incremental translation over growing audio chunks; apply a two-stage RL post-training on an offline ST model with rewards defined on translation behavior under partial audio, then deploy with adjustable chunk sizes (no separate policy). On ACL 60/60 and MCIF (EnDe/Zh/Ja), RL-adapted models achieve better quality–latency trade-offs than strong dialogue-style SimulST baselines, without trajectory supervision or explicit latency objectives.

Foundation Model Theory

The Emergence and Sparsification of a Syntactic Circuit [arxiv]

Qing Yao, Sasha Boguraev, Tiago Pimentel, Kyle Mahowald

A model can learn a syntactic skill early, yet the causal circuit implementing it keeps reorganizing and becomes much sparser later in training.

Tracks the subject–verb number agreement circuit across training (not just at convergence), showing a two-phase dynamic: early behavioral acquisition followed by prolonged internal reconfiguration and sparsification—so "the circuit for X" is not fixed. At many checkpoints, causally identify attention heads/MLP neurons mediating number information from subject to verb in Pythia-1B and across 10 seeds of Pythia-410M; compare stable task performance to changing internal pathways. Agreement accuracy emerges early and stabilizes, but the mediating components keep changing (e.g., dominant heads shift; in Pythia-1B a later head replaces early ones). Over time the mechanism contracts into a strikingly sparse late-training circuit, implying late training consolidates implementation even after capability plateaus.

LLM Evaluation

LLM-as-a-Judge Version Shift: The Instrument Changed, Not the Standard [arxiv]

Guijia Zhang, Harry Yang

Updating an LLM judge can change the evaluation instrument (criteria weighting and score resolution), not just shift scores—causing silent pipeline drift.

Preregistered study of judge version shifts showing changes in rubric geometry and discrimination power; proposes a low-cost "anchor audit" to detect offset vs reweighting vs resolution collapse before deploying a new judge. Hold judged items fixed (ML research proposals) and vary only judge version across three adjacent transitions (Qwen/Gemma/Mistral). Randomize information "packets" to separate formatting/presentation effects; fit a measurement view (origin shift, criterion rotation/warp, resolution change) and operationalize via

~200-call stratified anchor audit. Version changes aren't simply "harsher": two transitions are more lenient; one is inconsistent (anchors vs non-anchors). One transition collapses ranking power (93% of scores at one value), and identical facts don't remove gaps; packet tests implicate formatting sensitivity. Global recalibration fixes means but not item-level disagreement; the anchor audit flags real post-registration API judge swaps.

Field Overview

Today's arXiv batch is dominated by LLM agents and their infrastructure: at least ~40–60 papers touch agent workflows, long-horizon interaction, memory (e.g., SLICE-Mem/FlexA-Mem/Memory Tax/ActionRAG), and agent evaluation/harness design (e.g., -Bench, Trial by Trace, Async-RBench, Compact-Bench, SWE-Shift), with a noticeable subcluster on multi-agent coordination and failure modes (collusion/deception, communication failures, negotiation/ToM). A second major theme is reliability/safety/calibration for both base models and agentic RAG—roughly ~30+ papers on hallucination detection/mitigation, evidence sufficiency/abstention, jailbreaks and prompt-injection/tool-output attacks, and calibrated decision models (Jev-style, conformal factuality, confidence targets). Efficiency is the other big pillar (another ~25–40): long-context compression/KV-cache reuse, speculative decoding, quantization/FP4 training, sparsity/MoE compression, and looped/recurrent-depth Transformers (StateLoop/LoopFold/Accordion/halting heads). A somewhat surprising cluster is the breadth of audio/speech foundation work (20–30 papers) spanning ASR/TTS, full-duplex spoken agents, audio reasoning with RL, and voice security/deepfake/impersonation—suggesting "agentic + multimodal" is expanding beyond vision into speech as a first-class frontier.